

Problem 1:

$$p(D|\theta) = \prod_{i=1}^n p(x_i|\theta)$$

20

$$l(\theta) = \ln p(D|\theta)$$

$$\begin{aligned} \text{gradient}_{\theta}(l(\theta)) &= \sum_{i=1}^n \text{gradient}_{\theta}(\ln p(x_i|\theta)) \\ &= \sum_{i=1}^n \text{gradient}_{\theta}(\ln(\theta)) - \text{gradient}_{\theta}(\theta * x_i) \end{aligned}$$

$$0 = n / \theta_{\text{hat}} - \sum_{i=1}^n x_i$$

$$\text{sum}(x_i) = n / \theta_{\text{hat}}$$

$$\theta_{\text{hat}} = n / \text{sum}(x_i)$$

$$\theta_{\text{hat}} = 1 / (\text{sum}(x_i) * 1/n)$$

Averages for 10 runs each (code attached):

Max Likelihood Estimation

Average: 75.78%

Std. Dev: 2.302

Naive-Bayes

Average: 73.65%

Std. Dev: 2.409

k-Nearest Neighbor k = 1

Average: 64.92%

Std. Dev: 1.680

k-Nearest Neighbor k = 5

Average: 65.49%

Std. Dev: 1.100

k-Nearest Neighbor k = 11

Average: 64.97%

Std. Dev: 2.832

Parzen Window edge length = 20

Average: 65.03%

Std. Dev: 1.610

Problem 2: 25

Code is perfect, however, you are not coordinating the two classifiers so that they operate on same train/test data (-5). Also, there is no conclusion (-10).

Problem 3: 15

Accuracy should increase with k. This is due to bug. See later.

Problem 4: 12

Problem 2 (Maximum Likelihood Estimation):

```
import csv
import math
import numpy as np

runs = 0
accuracy = []
while (runs < 10):
    raw_data = np.genfromtxt("pima-indians-diabetes.csv", delimiter =
    ",", usecols = (1, 2, 3, 8))
    np.random.shuffle(raw_data) #randomize data

    train_data = raw_data[0:len(raw_data)/2] #assign 1st half of data
to training data
    test_data = raw_data[len(raw_data)/2:] #assign 2nd half of data
to test data
    correct = 0
    wrong = 0

    train_data_filter1 = train_data[:, 3] = train_data[:, 3] == 1
#setting filter for training data when column 9 = 1
    mean1 = train_data[train_data_filter1].mean(0) #calculate mean of
each attribute
    mean1 = np.delete(mean1, 3, 0) #remove the 4th column from mean
vector
    var1 = np.cov(train_data[train_data_filter1].transpose())
#calculate covariance matrix
    var1 = np.delete(var1, 3, 1) #removes the 4th column from the
covariance matrix
    var1 = np.delete(var1, 3, 0) #removes the 4th row from the
covariance matrix

    train_data_filter0 = train_data[:, 3] = train_data[:, 3] == 0
    mean0 = train_data[train_data_filter0].mean(0)
    mean0 = np.delete(mean0, 3, 0)
    var0 = np.cov(train_data[train_data_filter0].transpose())
    var0 = np.delete(var0, 3, 1)
    var0 = np.delete(var0, 3, 0)

    priortmp1 = train_data[train_data_filter1].shape[0] #gives the
count where training data column 4 equals 1
    priortmp0 = train_data[train_data_filter0].shape[0]

    prior0 = float(priortmp0) / (priortmp0 + priortmp1) #calculates
prior probability of trainging data where column 4 equals 0
    prior1 = float(priortmp1) / (priortmp0 + priortmp1)

    constant1 = 1 / math.sqrt(np.linalg.det(var1)) #calculates
1/sqrt(covariance)
    constant0 = 1 / math.sqrt(np.linalg.det(var0))
    inv_var0 = np.linalg.inv(var0) #inverse of covariance matrix
```

```

    inv_var1 = np.linalg.inv(var1)

    for i in test_data:
        j = np.delete(i, 3, 0)
        diff0 = j - mean0
        lklhood0 = constant0 * np.exp(-np.dot(np.dot(diff0,
inv_var0), diff0) / 2)

        diff1 = j - mean1
        lklhood1 = constant1 * np.exp(-np.dot(np.dot(diff1,
inv_var1), diff1) / 2)

        post0 = prior0 * lklhood0
        post1 = prior1 * lklhood1

        if(post0 > post1 and i[3] == 0):
            correct += 1
        elif(post1 < post0 and i[3] == 1):
            correct += 1
        else:
            wrong += 1

    accuracy.append(100 * float(correct) / (correct + wrong))
    runs += 1

print "Average Accuracy: " + str(np.average(accuracy))
print "Standard Deviation of Accuracy: " + str(np.std(accuracy))

```

Problem 2 (Naive Bayes):

```
import csv
import math
import numpy as np

runs = 0
accuracy = []
while(runs < 10):
    raw_data = np.genfromtxt("pima-indians-diabetes.csv", delimiter =
    ",", usecols = (1, 2, 3, 8))
    np.random.shuffle(raw_data) #randomize raw data

    train_data = raw_data[0:len(raw_data)/2] #separate 1st half to
training data
    test_data = raw_data[len(raw_data)/2:] #separate 2nd half to test
data
    correct = 0
    wrong = 0

    train_data_filter1 = train_data[:, 3] = train_data[:, 3] == 1
#setup filter where training data column 4 = 1
    mean1 = train_data[train_data_filter1].mean(0) #calculate mean
for each attribute in training data
    mean1 = np.delete(mean1, 3, 0) #delete the 4th column from mean
vector
    var1 = np.cov(train_data[train_data_filter1].transpose())
#calculate covariance matrix on training data
    var1 = np.delete(var1, 3, 1) #delete column 4 from covariance
matrix
    var1 = np.delete(var1, 3, 0) #delete row 4 from covariance matrix

    train_data_filter0 = train_data[:, 3] = train_data[:, 3] == 0
#setup filter where training data column 4 = 0
    mean0 = train_data[train_data_filter0].mean(0)
    mean0 = np.delete(mean0, 3, 0)
    var0 = np.cov(train_data[train_data_filter0].transpose())
    var0 = np.delete(var0, 3, 1)
    var0 = np.delete(var0, 3, 0)

    priortmp1 = train_data[train_data_filter1].shape[0] #number of
training data records where column 4 = 1
    priortmp0 = train_data[train_data_filter0].shape[0] #number of
training data records where column 4 = 0

    #calculate priors
    prior0 = float(priortmp0) / (priortmp0 + priortmp1)
    prior1 = float(priortmp1) / (priortmp0 + priortmp1)

    for i in test_data:
        #calculate the independent likelihoods
```

```

        lklhood00 = np.exp(-(i[0] - mean0[0]) * (i[0] - mean0[0]) /
(2 * var0[0,0])) / math.sqrt(var0[0,0])
        lklhood01 = np.exp(-(i[1] - mean0[1]) * (i[1] - mean0[1]) /
(2 * var0[1,1])) / math.sqrt(var0[1,1])
        lklhood02 = np.exp(-(i[2] - mean0[2]) * (i[2] - mean0[2]) /
(2 * var0[2,2])) / math.sqrt(var0[2,2])

        lklhood10 = np.exp(-(i[0] - mean1[0]) * (i[0] - mean1[0]) /
(2 * var1[0,0])) / math.sqrt(var1[0,0])
        lklhood11 = np.exp(-(i[1] - mean1[1]) * (i[1] - mean1[1]) /
(2 * var1[1,1])) / math.sqrt(var1[1,1])
        lklhood12 = np.exp(-(i[2] - mean1[2]) * (i[2] - mean1[2]) /
(2 * var1[2,2])) / math.sqrt(var1[2,2])

        #calculate the posteriors
        post0 = prior0 * lklhood00 * lklhood01 * lklhood02
        post1 = prior1 * lklhood10 * lklhood11 * lklhood12

        #make decisions based on the posterior and count correct vs
wrong on test data
        if(post0 > post1 and i[3] == 0):
            correct += 1
        elif(post1 < post0 and i[3] == 1):
            correct += 1
        else:
            wrong += 1

    runs += 1
    accuracy.append(100 * float(correct) / (correct + wrong))

print "Accuracy Average: " + str(np.average(accuracy))
print "Accuracy Standard Deviation: " + str(np.std(accuracy))

```

Problem 3 (k-Nearest Neighbor):

```
import csv
import math
import numpy as np
from scipy import spatial

runs = 0
accuracy = []
while(runs < 10):
    raw_data = np.genfromtxt("pima-indians-diabetes.csv", delimiter =
    ",", usecols = (1, 2, 3, 8))
    np.random.shuffle(raw_data) #shuffle data
    num_samples = 1 #value of k

    train_data = raw_data[0:len(raw_data)/2] #separate 1st half of
raw data into training data
    test_data = raw_data[len(raw_data)/2:] #separate 2nd half of raw
data into test data
    correct = 0
    wrong = 0
    post0 = 0
    post1 = 0

    train = np.delete(train_data, 3, 1) #assign train to train_data
without column 4
    tree = spatial.KDTree(train) #create kd tree with training data

    for i in test_data:
        distance, closest = tree.query(np.delete(i, 3, 0), k =
num_samples) #find k nearest neighbor(s) and assign it/them to closest
        j = 0
        #closest will be a list if k is greater than 1 and integer
if equal to 1, so this conditional separates them
        if(num_samples > 1):
            while(j < num_samples): #this checks each of the
closest neighbor
                if(train_data[closest[j], 3] == 1):
                    post1 += 1
                else:
                    post0 += 1
                j += 1
        else:
            if(train_data[closest, 3] == 1):
                post1 += 1
            else:
                post0 += 1

        if(post1 > post0 and i[3] == 1):
            correct += 1
        elif(post0 > post1 and i[3] == 0):
            correct += 1
```

```
        else:
            wrong += 1

    runs += 1
    accuracy.append(100 * float(correct) / (correct + wrong))

print "Accuracy Average: " + str(np.average(accuracy))
print "Accuracy Standard Deviation: " + str(np.std(accuracy))
```

Problem 4 (Parzen Window):

```
import csv
import math
import numpy as np
from scipy import spatial

runs = 0
accuracy = []
while(runs < 10):
    raw_data = np.genfromtxt("pima-indians-diabetes.csv", delimiter =
    ",", usecols = (1, 2, 3, 8))
    np.random.shuffle(raw_data) #randomizes data
    window_size = 20 #sets radius of hypercube

    train_data = raw_data[0:len(raw_data)/2] #separates the 1st half
of data to training data
    test_data = raw_data[len(raw_data)/2:] #separates the 2nd half of
data to testing data
    correct = 0
    wrong = 0
    post0 = 0
    post1 = 0

    train = np.delete(train_data, 3, 1) #assign training data without
column 4 to train
    tree = spatial.KDTree(train) #assign training data to kd tree

    for i in test_data:
        closest = tree.query_ball_point(np.delete(i, 3, 0),
window_size) #returns the indexes for neighbors inside of window
        j = 0
        while(j < len(closest)):
            if(train_data[closest[j], 3] == 1):
                post1 += 1
            else:
                post0 += 1
            j += 1

        if(post1 > post0 and i[3] == 1):
            correct += 1
        elif(post0 > post1 and i[3] == 0):
            correct += 1
        else:
            wrong += 1

    runs += 1
    accuracy.append(100 * float(correct) / (correct + wrong))

print "Accuracy Average: " + str(np.average(accuracy))
print "Accuracy Standard Deviation: " + str(np.std(accuracy))
```